

---

---

# Advanced Recommender Systems for E-commerce: A Comparative Study of Classical, Hybrid, and Neural Approaches

---

---

Final Project for the Course of  
ADVANCED MACHINE LEARNING  
Master's Program in Applied Artificial Intelligence  
By

ANGSAR SHAUMEN

Under the Supervision of  
(Zhanar Oralbekova)



SCHOOL OF ARTIFICIAL INTELLIGENCE AND DATA SCIENCE  
ASTANA IT UNIVERSITY LLP  
KAZAKHSTAN

FEBRUARY 2026

# Abstract

This work investigates the performance, robustness, and practical trade-offs of classical, hybrid, and neural recommender systems for e-commerce under realistic constraints such as data sparsity and cold-start scenarios.

We construct a unique dataset derived from Amazon Electronics reviews, specifically filtering for the Headphones subcategory and applying 5-core filtering to ensure meaningful interaction density. The task is formulated as implicit feedback Top-K recommendation with strict time-based train/validation/test splits to prevent data leakage.

Four representative models are implemented and compared: Popularity-based baseline, Item-KNN collaborative filtering, Matrix Factorization (MF), and Neural Collaborative Filtering (NeuMF/NCF). Evaluation is conducted using ranking metrics including Precision@K, Recall@K, and NDCG@K.

Our experiments reveal that NeuMF achieves the highest NDCG@10 (0.0059), outperforming both classical Matrix Factorization (0.0021) and the Popularity baseline (0.0050). The results demonstrate that neural architectures can effectively capture non-linear user-item interactions even in highly sparse settings (97.4% sparsity), though popularity remains a strong baseline due to the long-tail distribution of interactions.

**Keywords:** Recommender Systems; Collaborative Filtering; Neural Networks; E-commerce; Cold-Start Problem

# List of Tables

| Table                                                                  | Page |
|------------------------------------------------------------------------|------|
| 2.1 Dataset Statistics After Processing . . . . .                      | 5    |
| 2.2 Model Performance Comparison . . . . .                             | 8    |
| 2.3 Ablation Study: Impact of Embedding Dimension on NDCG@10 . . . . . | 10   |

# List of Figures

| Figure                                                                     | Page |
|----------------------------------------------------------------------------|------|
| 2.1 Long-tail distribution of item interactions (log-log scale) . . . . .  | 5    |
| 2.2 Model performance comparison across Precision, Recall, and NDCG@10 . . | 8    |
| 2.3 Training loss over epochs for MF and NCF . . . . .                     | 8    |
| 2.4 Cold-start distribution in the test set . . . . .                      | 9    |

# Table of Contents

|                                                        |            |
|--------------------------------------------------------|------------|
| <b>Abstract</b>                                        | <b>i</b>   |
| <b>List of Tables</b>                                  | <b>ii</b>  |
| <b>List of Figures</b>                                 | <b>iii</b> |
| <br>                                                   |            |
| <b>1 Introduction</b>                                  | <b>1</b>   |
| 1.1 Background and Motivation . . . . .                | 1          |
| 1.2 Related Work . . . . .                             | 1          |
| 1.3 Problem Statement . . . . .                        | 2          |
| 1.4 Research Questions . . . . .                       | 3          |
| 1.5 Contributions . . . . .                            | 3          |
| <br>                                                   |            |
| <b>2 Methodology and Experiments</b>                   | <b>4</b>   |
| 2.1 Dataset Description . . . . .                      | 4          |
| 2.1.1 Data Processing . . . . .                        | 4          |
| 2.1.2 Dataset Statistics . . . . .                     | 4          |
| 2.1.3 Data Distribution Analysis . . . . .             | 5          |
| 2.2 Problem Formulation . . . . .                      | 5          |
| 2.3 Models . . . . .                                   | 6          |
| 2.3.1 Popularity Baseline . . . . .                    | 6          |
| 2.3.2 Item-Based Collaborative Filtering . . . . .     | 6          |
| 2.3.3 Matrix Factorization . . . . .                   | 6          |
| 2.3.4 Neural Collaborative Filtering (NeuMF) . . . . . | 6          |
| 2.4 Evaluation Metrics . . . . .                       | 7          |
| 2.5 Experimental Setup . . . . .                       | 7          |
| 2.6 Results . . . . .                                  | 7          |

|          |                                               |           |
|----------|-----------------------------------------------|-----------|
| 2.6.1    | Main Results . . . . .                        | 7         |
| 2.6.2    | Training Dynamics . . . . .                   | 8         |
| 2.6.3    | Cold-Start Analysis . . . . .                 | 9         |
| 2.7      | Discussion . . . . .                          | 9         |
| 2.7.1    | Model Design Choices . . . . .                | 9         |
| 2.7.2    | Ablation Study: Embedding Dimension . . . . . | 10        |
| <b>3</b> | <b>Conclusion</b>                             | <b>11</b> |
| 3.1      | Summary of Findings . . . . .                 | 11        |
| 3.2      | Limitations . . . . .                         | 12        |
| 3.3      | Future Work . . . . .                         | 12        |
| 3.4      | Reproducibility . . . . .                     | 12        |

# Chapter 1

## Introduction

### 1.1 Background and Motivation

Recommender systems have become essential components of modern e-commerce platforms, helping users navigate vast product catalogs and discover relevant items [? ]. Major platforms such as Amazon, Netflix, and Spotify attribute significant portions of their revenue to recommendation-driven engagement.

The fundamental challenge in recommendation lies in learning user preferences from sparse interaction data. Traditional approaches based on collaborative filtering, such as Matrix Factorization, have proven effective by decomposing the user-item interaction matrix into latent factors [? ]. However, these linear models may fail to capture the complex, non-linear patterns inherent in user behavior.

Recent advances in deep learning have led to the development of Neural Collaborative Filtering (NCF), which replaces the inner product of latent factors with neural networks capable of learning arbitrary functions [? ]. This approach promises better generalization, particularly in sparse data regimes.

### 1.2 Related Work

The field of recommender systems has evolved significantly over the past two decades. This section surveys the key approaches relevant to our comparative study.

**Collaborative Filtering.** The foundational work by Sarwar et al. [? ] introduced

item-based collaborative filtering, demonstrating that computing similarities between items rather than users provides both computational efficiency and recommendation quality. This approach remains widely used as a baseline in academic research.

**Matrix Factorization.** The landmark work by Koren et al. [?] established matrix factorization as the dominant paradigm for collaborative filtering. Their SVD++ model extended basic MF with implicit feedback and temporal dynamics, achieving state-of-the-art performance on the Netflix Prize dataset.

**Implicit Feedback.** Hu et al. [?] pioneered methods for handling implicit feedback, recognizing that user behavior data (clicks, purchases) differs fundamentally from explicit ratings. Their weighted matrix factorization approach assigns varying confidence levels to observed and unobserved interactions.

**Learning to Rank.** Rendle et al. [?] proposed Bayesian Personalized Ranking (BPR), a pairwise learning framework optimized directly for ranking rather than prediction. BPR has become the de facto training objective for many modern recommender systems.

**Neural Approaches.** He et al. [?] introduced Neural Collaborative Filtering (NCF), replacing the inner product in MF with a multi-layer perceptron. Their work demonstrated that neural networks can learn more expressive user-item interaction functions, particularly beneficial under sparsity.

**Hybrid Methods.** The RecSys literature extensively covers hybrid approaches that combine collaborative and content-based signals. The LightFM library provides a practical implementation that learns unified embeddings for users, items, and features. However, hybrid methods require rich item metadata, which is not always available in implicit feedback datasets.

Our work contributes to this literature by providing a systematic comparison across paradigms using a carefully controlled evaluation framework with proper time-based splitting and ranking metrics.

## 1.3 Problem Statement

This research addresses the following question: **How do classical, hybrid, and neural recommender systems compare in terms of ranking quality, computational**



**efficiency, and robustness to cold-start scenarios?**

Specifically, we focus on:

1. Constructing a unique, reproducible dataset from public Amazon reviews
2. Implementing representative models from each paradigm (classical, neural)
3. Conducting rigorous evaluation using ranking metrics
4. Analyzing cold-start behavior and computational trade-offs

## **1.4 Research Questions**

1. How do classical, hybrid, and neural recommenders compare on ranking metrics?
2. Does the neural approach (NCF) consistently outperform Matrix Factorization?
3. How do models behave under extreme sparsity conditions?
4. What are the computational trade-offs between approaches?

## **1.5 Contributions**

The main contributions of this work are:

- A unique dataset derived from Amazon Electronics (Headphones subcategory) with 5-core filtering
- Implementation of four representative recommendation models
- Comprehensive experimental comparison using Precision, Recall, and NDCG metrics
- Analysis of cold-start scenarios and model robustness
- Reproducible codebase and experimental framework

# Chapter 2

## Methodology and Experiments

### 2.1 Dataset Description

We utilize a subset of the Amazon Product Reviews dataset, specifically focusing on the Electronics category with filtering for Headphones-related products. The dataset is sourced from the Stanford SNAP repository [? ].

#### 2.1.1 Data Processing

The following preprocessing steps ensure dataset uniqueness and research reproducibility:

1. **Subcategory Filtering:** Only products categorized under Headphones, Earbuds, or Over-Ear Headphones are retained
2. **Implicit Feedback Conversion:** All ratings are converted to binary implicit feedback (1 = interacted)
3. **5-Core Filtering:** Users and items with fewer than 5 interactions are iteratively removed
4. **Time-Based Split:** Data is split chronologically (70% train, 15% validation, 15% test)

#### 2.1.2 Dataset Statistics

Table 2.1 summarizes the final dataset characteristics after preprocessing.

Table 2.1: Dataset Statistics After Processing

| Metric              | Value  |
|---------------------|--------|
| Users               | 368    |
| Items               | 282    |
| Interactions        | 2,695  |
| Sparsity            | 97.40% |
| Train Set Size      | 1,888  |
| Validation Set Size | 407    |
| Test Set Size       | 400    |

### 2.1.3 Data Distribution Analysis

Figure 2.1 illustrates the long-tail distribution commonly observed in recommender systems, where a small number of items receive the majority of interactions.




Figure 2.1: Long-tail distribution of item interactions (log-log scale)

## 2.2 Problem Formulation

We formulate the recommendation task as follows:

- **Task:** Top-K Item Recommendation
- **Feedback Type:** Implicit (binary interactions)
- **Output:** Ranked list of items for each user
- **Evaluation:** Ranking metrics (Precision, Recall, NDCG)

## 2.3 Models

### 2.3.1 Popularity Baseline

The simplest baseline recommends items based on global popularity. For user  $u$ , the recommendation list consists of the  $K$  most frequently interacted items not already seen by  $u$ .

### 2.3.2 Item-Based Collaborative Filtering

Item-KNN computes similarity between items based on co-occurrence patterns. Given item  $i$  and item  $j$ , cosine similarity is computed as:

$$\text{sim}(i, j) = \frac{\mathbf{r}_i \cdot \mathbf{r}_j}{\|\mathbf{r}_i\| \cdot \|\mathbf{r}_j\|} \quad (2.1)$$

### 2.3.3 Matrix Factorization

Matrix Factorization decomposes the user-item interaction matrix  $R$  into low-rank matrices:

$$\hat{R} = P \cdot Q^T \quad (2.2)$$

where  $P \in \mathbb{R}^{|U| \times d}$  and  $Q \in \mathbb{R}^{|I| \times d}$  are user and item embedding matrices.

### 2.3.4 Neural Collaborative Filtering (NeuMF)

NeuMF combines Generalized Matrix Factorization (GMF) with a Multi-Layer Perceptron (MLP) [? ]:

$$\hat{y}_{ui} = \sigma(h^T(\phi_{GMF} \oplus \phi_{MLP})) \quad (2.3)$$

The architecture consists of:

- User and item embeddings for GMF path
- User and item embeddings for MLP path
- Hidden layers with ReLU activation
- Final prediction layer with sigmoid output

## 2.4 Evaluation Metrics

We evaluate using standard ranking metrics:

**Precision@K:** The fraction of recommended items that are relevant

$$Precision@K = \frac{|Rec_K \cap Relevant|}{K} \quad (2.4)$$

**Recall@K:** The fraction of relevant items that are recommended

$$Recall@K = \frac{|Rec_K \cap Relevant|}{|Relevant|} \quad (2.5)$$

**NDCG@K:** Normalized Discounted Cumulative Gain

$$NDCG@K = \frac{DCG@K}{IDCG@K} \quad (2.6)$$

## 2.5 Experimental Setup

All experiments are conducted with the following configuration:

- Embedding dimension: 32
- Learning rate: 0.001 (NCF), 0.005 (MF)
- Batch size: 256
- Negative samples: 4 per positive
- Training epochs: 30
- Hardware: NVIDIA RTX 2060 GPU

## 2.6 Results

### 2.6.1 Main Results

Table 2.2 presents the main experimental results across all models.

Table 2.2: Model Performance Comparison

| Model              | P@5           | P@10          | R@5           | R@10          | NDCG@5        | NDCG@10       |
|--------------------|---------------|---------------|---------------|---------------|---------------|---------------|
| Popularity         | 0.0023        | 0.0023        | 0.0043        | 0.0100        | 0.0029        | 0.0050        |
| Item-KNN           | 0.0000        | 0.0006        | 0.0000        | 0.0057        | 0.0000        | 0.0020        |
| MF                 | 0.0011        | 0.0017        | 0.0011        | 0.0054        | 0.0008        | 0.0021        |
| <b>NeuMF (NCF)</b> | <b>0.0034</b> | <b>0.0023</b> | <b>0.0071</b> | <b>0.0100</b> | <b>0.0042</b> | <b>0.0059</b> |

Figure 2.2 visualizes the model comparison across metrics.




Figure 2.2: Model performance comparison across Precision, Recall, and NDCG@10

## 2.6.2 Training Dynamics

Figure 2.3 shows the training loss curves for MF and NCF models.




Figure 2.3: Training loss over epochs for MF and NCF

### 2.6.3 Cold-Start Analysis

A significant portion of test users (48.57%) and items (51.28%) are "cold" (not seen in training), posing challenges for collaborative filtering approaches.




Figure 2.4: Cold-start distribution in the test set

## 2.7 Discussion

**Key Finding 1:** NeuMF achieves the highest NDCG@10 (0.0059), demonstrating that neural architectures can capture complex interaction patterns even in highly sparse settings.

**Key Finding 2:** The Popularity baseline remains competitive (NDCG@10 = 0.0050), reflecting the strong long-tail effect where popular items dominate.

**Key Finding 3:** Classical Matrix Factorization underperforms (NDCG@10 = 0.0021), likely due to the extreme sparsity making it difficult to learn meaningful latent factors.

**Key Finding 4:** Item-KNN shows near-zero performance on Precision@5, indicating that neighborhood-based methods struggle when the interaction matrix is extremely sparse.

### 2.7.1 Model Design Choices

**Why not SVD++?** While SVD++ extends basic Matrix Factorization with user and item biases, our experiments with extreme sparsity (97.4%) suggest that bias terms provide minimal benefit. The limited interaction data makes it difficult to reliably estimate

per-user and per-item biases. We therefore opted for the simpler MF formulation to avoid overfitting.

**Why not LightFM (Hybrid)?** The LightFM library enables hybrid recommendations by incorporating item and user features alongside collaborative signals. However, the Amazon Reviews dataset in our filtered subcategory lacks rich item metadata (e.g., detailed product attributes, brand information). Without meaningful content features, hybrid models reduce to pure collaborative filtering, negating their primary advantage. We acknowledge this as a limitation and suggest that future work with richer metadata could benefit from hybrid approaches.

### 2.7.2 Ablation Study: Embedding Dimension

We conducted an ablation study to analyze the impact of embedding dimension on model performance. Table 2.3 presents the results for embedding dimensions of 16, 32, 64, and 128.

Table 2.3: Ablation Study: Impact of Embedding Dimension on NDCG@10

| Embedding Dim | MF     | NCF           |
|---------------|--------|---------------|
| 16            | 0.0063 | 0.0069        |
| 32            | 0.0073 | <b>0.0119</b> |
| 64            | 0.0057 | 0.0089        |
| 128           | 0.0010 | 0.0120        |

The results indicate that an embedding dimension of 32 provides optimal performance for both models. Lower dimensions (16) lack sufficient representational capacity, while higher dimensions (64, 128) lead to overfitting on the sparse interaction data. This finding is consistent with prior literature suggesting that moderate embedding sizes work best under sparsity constraints.



# Chapter 3

## Conclusion

### 3.1 Summary of Findings

This work presented a comprehensive comparative study of recommender system approaches for e-commerce, specifically focusing on the Headphones subcategory of Amazon Electronics reviews.

Our main findings are:

1. **Neural methods show promise:** NeuMF (NCF) achieved the highest ranking performance ( $\text{NDCG@10} = 0.0059$ ), outperforming both classical Matrix Factorization and baseline approaches.
2. **Sparsity remains challenging:** With 97.4% sparsity, traditional collaborative filtering methods like Item-KNN and MF struggled to learn meaningful patterns.
3. **Popularity is a strong baseline:** The simple popularity-based recommender achieved competitive results ( $\text{NDCG@10} = 0.0050$ ), highlighting the importance of the long-tail distribution in e-commerce data.
4. **Cold-start is significant:** Approximately 50% of test users and items were not seen during training, underscoring the importance of cold-start handling in production systems.

## 3.2 Limitations

This study has several limitations:

- The dataset size after filtering is relatively small (2,695 interactions), which may limit the generalizability of findings
- Hybrid models incorporating content features (e.g., LightFM) were not fully evaluated due to technical constraints
- The study focuses on offline evaluation; online A/B testing would provide stronger evidence of practical utility

## 3.3 Future Work

Future research directions include:

1. **Hybrid approaches:** Incorporating item metadata (category, brand, price) to address cold-start
2. **Sequential models:** Exploring transformer-based architectures (e.g., SASRec, BERT4Rec) for session-based recommendation
3. **Larger datasets:** Evaluating on larger subcategories or the full Amazon dataset
4. **Explainability:** Investigating interpretable recommendation methods

## 3.4 Reproducibility

All code, data processing scripts, and trained models are available in the accompanying repository. The experiments can be reproduced using the provided `run_all.sh` script and notebook files.

# Bibliography

- [1] Francesco Ricci, Lior Rokach, and Bracha Shapira.  
Introduction to recommender systems handbook.  
*Recommender Systems Handbook*, pages 1–35, 2011.
- [2] Yehuda Koren, Robert Bell, and Chris Volinsky.  
Matrix factorization techniques for recommender systems.  
*Computer*, 42(8):30–37, 2009.
- [3] Xiangnan He, Lizi Liao, Hanwang Zhang, Liqiang Nie, Xia Hu, and Tat-Seng Chua.  
Neural collaborative filtering.  
In *Proceedings of the 26th International Conference on World Wide Web*, pages 173–182, 2017.
- [4] Badrul Sarwar, George Karypis, Joseph Konstan, and John Riedl.  
Item-based collaborative filtering recommendation algorithms.  
*Proceedings of the 10th International Conference on World Wide Web*, pages 285–295, 2001.
- [5] Yifan Hu, Yehuda Koren, and Chris Volinsky.  
Collaborative filtering for implicit feedback datasets.  
In *2008 Eighth IEEE International Conference on Data Mining*, pages 263–272, 2008.
- [6] Steffen Rendle, Christoph Freudenthaler, Zeno Gantner, and Lars Schmidt-Thieme.  
Bpr: Bayesian personalized ranking from implicit feedback.  
In *Proceedings of the 25th Conference on Uncertainty in Artificial Intelligence*, pages 452–461, 2009.
- [7] Julian McAuley, Christopher Targett, Qinfeng Shi, and Anton Van Den Hengel.  
Image-based recommendations on styles and substitutes.  
In *Proceedings of the 38th International ACM SIGIR Conference*, pages 43–52, 2015.